


# Elasticsearch

БГУ ФПМИ, 2017


# Elasticsearch

- ▶ open-source distributed search server
- ▶ implementation is in Java
- ▶ powered by **Apache Lucene**
- ▶ to achieve fast search responses, instead of searching the text directly, it searches an **index** instead.

## Indexing:




## Data representation:



## Query DSL:

### basic queries



### compound queries



## main query structure

```
curl -X POST "http://localhost:9200/blog/_search?pretty=true" -d '{
  "from": 0,
  "size": 10,
  "query" : QUERY_JSON,
    FILTER_JSON,
    FACET_JSON,
    SORT_JSON
  }'
```

# Structure

`http://hostname:port/index_name/doc_type/doc_id`

- ▶ An **index** consists of one or more **documents**, and a **document** consists of one or more **fields**.
- ▶ A **type** is a logical category/partition of your index whose semantics is completely up to you.

# Concepts

## Near Realtime

There is a slight latency (normally one second) from the time you index a document until the time it becomes searchable.

## Shards & Replicas

Each index can be split into multiple shards (default=5). An index can also be replicated zero or more times (default=1).

# Install

```
wget https://artifacts.elastic.co/downloads/elasticsearch/  
elasticsearch-5.2.1.tar.gz
```

```
tar -zxvf elasticsearch-5.2.1.tar.gz
```

```
elasticsearch-5.2.1/bin/elasticsearch
```

```
sudo pip3 install elasticsearch
```

## Run

```
./elasticsearch -Ecluster.name=c_name -Enode.name=n_name
```

```
http://localhost:9200/
```

```
{
  "name" : "XHJP8FD",
  "cluster_name" : "elasticsearch",
  "cluster_uuid" : "LbmdPlaRTmKVbULaI6fmUg",
  "version" : {
    "number" : "5.2.1",
    "build_hash" : "db0d481",
    "build_date" : "2017-02-09T22:05:32.386Z",
    "build_snapshot" : false,
    "lucene_version" : "6.4.1"
  },
  "tagline" : "You Know, for Search"
}
```

# Search Basics



## Create Index

```
curl -XPUT 'localhost:9200/fox_ind' -d '{  
  "settings": {  
    "index.number_of_replicas": 0,  
    "index.number_of_shards": 1  
  }  
}'
```

## Let's Index!

```
curl -XPUT http://localhost:9200/fox_ind/texts/1 -d '{
  "content": "The quick brown fox"
}'
curl -XPUT http://localhost:9200/fox_ind/texts/2 -d '{
  "content": "Jumped over the lazy dog"
}'
curl -XPUT http://localhost:9200/fox_ind/texts/3 -d '{
  "content": "What does the fox say?"
}'
curl -XPUT http://localhost:9200/fox_ind/texts/4 -d '{
  "content": "The quick lazy brown fox did not jump."
}'
```

## Index in batch

```
curl -XPOST http://localhost:9200/fox_ind/texts/_bulk?pretty -d
'{"index": {"_id": "1"}}
{"content": "The quick brown fox"}
{"index": {"_id": "2"}}
{"content": "Jumped over the lazy dog"}
{"index": {"_id": "3"}}
{"content": "What does the fox say?"}
{"index": {"_id": "4"}}
{"content": "The quick lazy brown fox did not jump."}'
```

## GET /\_cat/

```
curl -XGET 'localhost:9200/_cat/health?v'
```

```
curl -XGET 'localhost:9200/_cat/nodes?v'
```

```
curl -XGET 'localhost:9200/_cat/indices?v'
```

```
curl -XGET 'localhost:9200/_cat/shards?v'
```

## GET /\_stats/

```
curl -XGET 'localhost:9200/fox_ind/_stats/docs?pretty=true'
```

```
curl -XGET 'localhost:9200/fox_ind/_stats/store?pretty=true  
&human=true'
```

## Shrink Index

```
curl -XPUT 'localhost:9200/fox_ind/_settings' -d '{
  "settings": {
    "index.blocks.write": true
  }}'
```

```
curl -XPOST 'localhost:9200/fox_ind/_shrink/fox_ind_shr' -d '{
  "settings": {
    "index.number_of_replicas": 0,
    "index.number_of_shards": 1
  }}'
```

## GET Document

Realtime operation:

```
curl -XGET 'localhost:9200/fox_ind/texts/1?pretty'
```

```
{
  "_index" : "fox_ind",
  "_type" : "texts",
  "_id" : "1",
  "_version" : 1,
  "found" : true,
  "_source" : {
    "content" : "The quick brown fox"
  }
}
```

## Simple Search

```
curl -XGET 'localhost:9200/fox_ind/_search?q=fox&pretty=true'
```

```
{
  "took" : 13,
  "timed_out" : false,
  "_shards" : {
    "total" : 5,
    "successful" : 5,
    "failed" : 0
  },
  "hits" : {
    "total" : 3,
    "max_score" : 0.561078,
    "hits" : [
      {
        "_index" : "fox_ind",
        "_type" : "texts",
        "_id" : "4",
        "_score" : 0.561078,
        "_source" : {
          "content" : "The quick lazy brown fox did not jump."
        }
      },
      ...
    ]
  }
}
```



## Boolean Search

"fox" AND "dog"

```
curl -XGET 'localhost:9200/fox_ind/_search?pretty' -d '{
  "query": {
    "bool": {
      "must": [
        {"match": {"content": "dog"} },
        {"match": {"content": "fox"} }
      ]
    }
  }
}'
```

"fox" OR "dog"

```
curl -XGET 'localhost:9200/fox_ind/_search?pretty' -d '{
  "query": {
    "bool": {
      "should": [
        {"match": {"content": "dog"} },
        {"match": {"content": "fox"} }
      ]
    }
  }
}'
```

## Fuzzy Match

```
curl -XPOST 'localhost:9200/fox_ind/_search?pretty' -d '{
  "query": {
    "match": {
      "content": {"query": "lary fix", "fuzziness": 1,
                  "operator": "AND"}
    }
  }
}'
```

"fuzziness": **"AUTO"** is recommended (fuzziness depends on term length).

# TF-IDF

## Term vectors

Для конкретного документа:

```
curl -XGET 'localhost:9200/fox_ind/texts/1/_termvectors?pretty' -d '{
  "fields": ["content"],
  "offsets": false,
  "payloads": false,
  "positions": false,
  "term_statistics": true,
  "field_statistics": true
}'
```

Статистика для термина по шарду:

```
curl -XGET 'localhost:9200/fox_ind/texts/_termvectors?pretty' -d '{
  "doc": {"content": "brown"},
  "term_statistics" : true, "positions" : false, "offsets" : false}'
```

## Term vectors: top-3 keywords

```
curl -XGET 'localhost:9200/fox_ind/texts/4/_termvectors?pretty'  
-d '{  
  "fields": ["content"],  
  "term_statistics" : true,  
  "offsets": false,  
  "positions": false,  
  "field_statistics": false,  
  "filter" : {"max_num_terms" : 3}  
}'
```

## Understanding `_score`

Lucene (and thus Elasticsearch) uses the Boolean model to find matching documents, and a formula called the **practical scoring function** to calculate relevance. This formula borrows concepts from tf-idf and the vector space model.

```
curl -XGET 'localhost:9200/fox_ind/texts/_search?q=fox&pretty&explain'
```

В версии 5.2.1 увидите стандартный BM25 (по-умолчанию без boost-добавок).

## Changing defaults

```
curl -XPOST 'localhost:9200/fox_ind/_close'
```

```
curl -XPUT 'localhost:9200/fox_ind/_settings?pretty' -d '{
  "settings": {
    "index": {
      "similarity": { "default": {
        "type": "BM25",
        "b": 0,
        "k1": 2
      }}
    }
  }
}'
```

```
curl -XPOST 'localhost:9200/fox_ind/_open'
```

## Другие similarity

- ▶ BM25
- ▶ Divergence from random
- ▶ TF-IDF
- ▶ ...

`https://www.elastic.co/guide/en/elasticsearch/  
reference/master/index-modules-similarity.html#  
\_available\_similarities`



# Analyzer

```
curl -XGET 'localhost:9200/fox_ind/_analyze?pretty' -d '{  
  "analyzer": "english", "text": "Text was semi-analyzed"}'
```

# Analyzer

```
{
  "settings": {
    "analysis": {
      "filter": {
        "russian_stop": {
          "type": "stop",
          "stopwords": "_russian_"
        },
        "russian_stemmer": {
          "type": "stemmer",
          "language": "russian"
        }
      },
      "analyzer": {
        "russian": {
          "tokenizer": "standard",
          "filter": [
            "lowercase",
            "russian_stop",
```